

# CS 2110 Homework 5

## LC-M

Madisson Grams, Vivian Thiebaut, Sam Gilson, Christopher Messina

Fall 2019

### Contents

<b>1</b>	<b>Overview</b>	<b>2</b>
<b>2</b>	<b>Instructions</b>	<b>2</b>
2.1	Part 1 . . . . .	2
2.1.1	Program Counter (PC) . . . . .	3
2.1.2	Instruction Register (IR) . . . . .	3
2.1.3	ADDR . . . . .	4
2.1.4	ALU . . . . .	4
2.1.5	SRC2MUX . . . . .	5
2.1.6	Condition Codes (CC) . . . . .	5
2.2	Part 2 . . . . .	5
2.2.1	Machine Code . . . . .	5
2.2.2	Instruction Format . . . . .	6
2.2.3	Loading the instructions to the LC-M . . . . .	7
2.3	Part 3 . . . . .	7
2.3.1	Information about State Machines . . . . .	7
<b>3</b>	<b>Testing your work</b>	<b>8</b>
<b>4</b>	<b>Deliverables</b>	<b>8</b>
<b>5</b>	<b>Demos</b>	<b>8</b>
<b>6</b>	<b>Rules and Regulations</b>	<b>9</b>
6.1	General Rules . . . . .	9
6.2	Submission Conventions . . . . .	9
6.3	Submission Guidelines . . . . .	9
6.4	Syllabus Excerpt on Academic Misconduct . . . . .	10
6.5	Is collaboration allowed? . . . . .	10

# 1 Overview

So far in CS 2110 you've learned the basics of digital logic and how to build simple circuits, but how does all of this become relevant to computer science? This assignment will help pull it altogether as you will complete missing circuits of a functional computer! This computer will be able to run simple programs and be a stepping stone into the rest of the material in this class.

## 2 Instructions

Your friend was building his own computer in CircuitSim, but then he accidentally deleted some portions of it. Now, your friend needs your help to fill in the missing parts! The computer you will be building is largely a subset of the LC-3 known as the LC-M. The LC-M supports only the following instructions:

- ADD
- AND
- BR
- LDR
- LEA
- NOT
- STR
- HALT

### 2.1 Part 1

We have provided `LCM.sim` that has five missing subcircuits:

- ALU
- CC-Logic
- PC
- ADDR
- SRC2MUX

Your job is to implement the missing parts of the LC-M. You **should not edit anything in the main circuit or other subcircuits**, however, feel free to take a look at them to get a better understanding of the rest of a computer!

Before starting, we recommend that you have a basic understanding of the LC-M's bigger sibling, the LC-3. Chapters 4 and 5 of the Patt Patel textbook are a good resource, as well as Appendix A, which goes into more detail about each instruction. Appendix C describes the behavior of the State Machine circuit in orchestrating the control signals. The LC-M state machine, while it functions similarly to the LC-3, has 32 states, uses a 28-bit control store for the control signals and uses two other stores for choosing the next state. The only different machine instruction is HALT which indicates the end of a program and replaces the LC-3 TRAP instruction. In the LC-3, TRAP allows for calling a variety of other pre-created instructions in memory. You do not have to worry about TRAPs for this assignment and you should merely interpret the TRAP opcode as HALT.

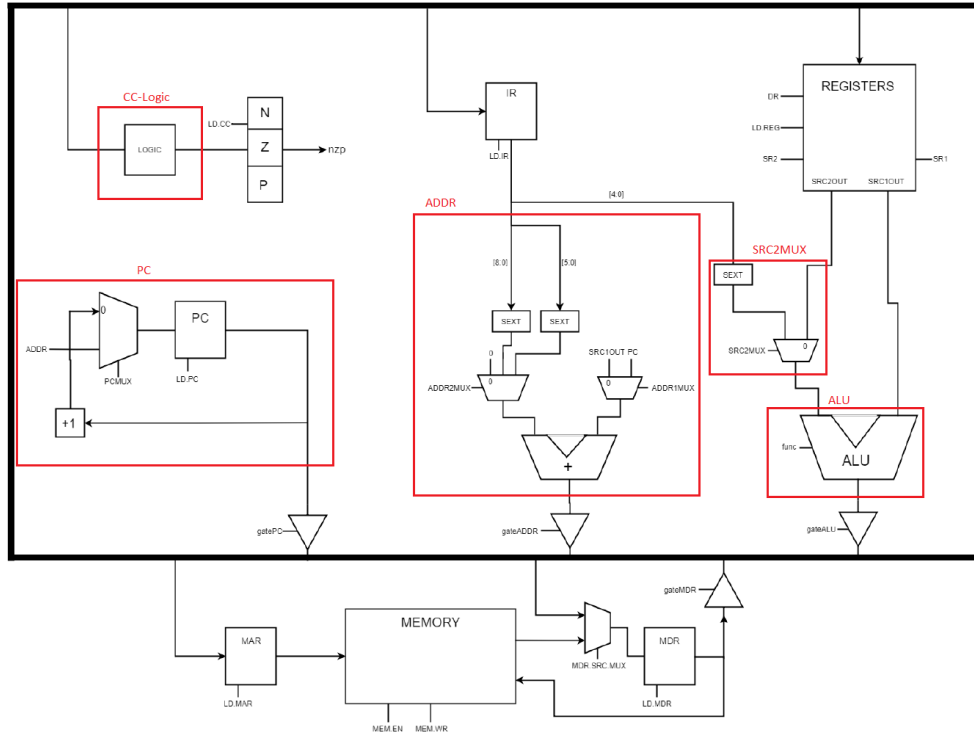


Figure 1: Diagram of the LC-M

### 2.1.1 Program Counter (PC)

The PC is a 16 bit register that holds the address of the next instruction to be executed. During the FETCH stage, the contents of the PC are loaded into the memory address register (MAR), and the PC is updated with the address of the next instruction. There are two scenarios for updating the PC:

1. The contents of the PC are incremented by 1.  
**Selected when PCMUX = 0**
2. The result of the ADDR is the address of the next instruction. The output from the ADDR should be stored in the PC. This occurs if we use the branching instruction, (BR).  
**Selected when PCMUX = 1**

**Implement this circuit in the subcircuit PC**

### 2.1.2 Instruction Register (IR)

The IR is a 16 bit register that holds the machine code of the current instruction. During the FETCH stage, the LC-M interrogates memory for the instruction, which gets loaded into the memory data register (MDR). Then, the contents of the MDR are loaded into the IR. The contents of the IR are then decoded and executed.

### Sign Extension

There are two LC-M instructions that use constants: ADD and AND. For example, ADD R0, R1, 2 adds the constant 2 to the contents of register 1 and stores the result in register 0. This constant is called imm5 (immediate 5) and is stored in the 5 least significant bits of the instruction. These 5 bits are sign extended to 16 bits so that imm5 can be used in other components of the LC-M.

### 2.1.3 ADDR

The ADDR is used to calculate addresses during the following instructions:

- BR | 0000 | nzp | PCoffset9 |

Here, the PCoffset9 (the 9 least significant bits of the instruction) is sign extended to 16 bits and added to the contents of the PC.

- LDR | 0110 | DR | baseR | offset6 |

Offset6 (the 6 least significant bits of the instruction) is sign extended to 16 bits and added to the contents of the base register, which can be accessed by SRC1OUT.

- LEA | 1110 | DR | PCoffset9 |

PCoffset9 (the 9 least significant bits of the instruction) is sign extended to 16 bits and added to the contents of the PC.

- STR | 0111 | SR | baseR | offset6 |

Offset6 (the 6 least significant bits of the instruction) is sign extended to 16 bits and added to the contents of the base register, which can be accessed by SRC1OUT.

Notice that in each of these instructions we are adding either SRC1OUT or PC to either sign-extended offset6 or sign-extended offset9 from IR. You will need two multiplexers for this part. To match up with our grader, ensure the following select bits match up with the inputs:

- 0: ADDR2MUX = 00
- SEXT(IR[8:0]): ADDR2MUX = 01
- SEXT(IR[5:0]): ADDR2MUX = 10
- SRC1OUT: ADDR1MUX = 0
- PC: ADDR1MUX = 1

**Implement this circuit in the ADDR subcircuit.**

### 2.1.4 ALU

The LC-M uses an ALU to execute various instructions, such as ADD and NOT. The ALU takes in two 16 bit inputs and has four operations:

0.  $A + B$
1.  $A \text{ AND } B$
2. NOT A
3. PASS A (A is unchanged)

**Implement this circuit in the ALU subcircuit**

### 2.1.5 SRC2MUX

The LC-M's ALU has two inputs: the A input is from SRC1OUT and the B input is from SRC2MUX. The SRC2MUX is a one-bit multiplexer that selects the second input to the ALU. It has two choices depending on the value of bit 5 in the IR:

1. SRC2OUT. This refers to the contents of the second source register in the instruction. For example, for the instruction `AND R0, R1, R2`, the register file outputs the contents of R2 as SRC2OUT and SRC2MUX selects this as the B input to the ALU.

**Selected when SRC2MUX = 0**

2. imm5. imm5 is the 5-bit constant in the instruction word. For example, the instruction `AND R0, R1, #4`, SRC2MUX selects the sign-extended 5-bit imm5 from the bits 0-4 of the instruction.

**Selected when SRC2MUX = 1**

**Implement this circuit in the SRC2MUX subcircuit**

### 2.1.6 Condition Codes (CC)

The LC-M has three condition codes: N (negative), Z (zero), and P (positive). These codes are set after the LC-M executes instructions that include loading a result into a register, such as `ADD` and `LDR`. For example, if `ADD R2, R0, R1` results in a positive number, then `NZP = 001`. Opcodes that are followed by “+” in Figure 2 are operations that set the condition codes. The CC subcircuit should set N to 1 if the input is negative, set Z to 1 if the input is zero, and set P to 1 if the input is positive. Only 1 bit in NZP should be set at any given time. Zero is not considered a positive number.

With that in mind, set the correct bit and **implement this circuit in the CC-Logic subcircuit**

Hint: you are allowed to use a comparator, which can be found under Arithmetic!

## 2.2 Part 2

Now that we've finished building the LC-M, we want to run a program with it! We have provided a simple assembly program which you will need to translate to machine code so that you can execute it on the LC-M. The program is equivalent to:

```
A = 15;
for (int i = 0; i < 4; i++) {
    A = A - i;
}
```

### 2.2.1 Machine Code

So far in most Computer Science courses, you've likely mostly only been exposed to high-level languages such as Python or Java and while these are readable somewhat easily by a human, they cannot be directly interpreted by a computer. The language of computers is binary. The process to get to binary for a language like C (which you will learn later in this course!) is compilation to assembly and then assembling that to machine code which is written in binary. You will function as the assembler for this program.

### 2.2.2 Instruction Format

The LC-M is set up to interpret instructions in a certain format. The chart below shows how instructions are stored to help you understand how to convert the assembly instructions to binary.

	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
ADD <sup>+</sup>	0001				DR			SR1			0	00		SR2		
ADD <sup>+</sup>	0001				DR			SR1			1	imm5				
AND <sup>+</sup>	0101				DR			SR1			0	00		SR2		
AND <sup>+</sup>	0101				DR			SR1			1	imm5				
BR	0000				n	z	p	PCOffset9								
LDR <sup>+</sup>	0110				DR			BaseR		offset6						
LEA <sup>+</sup>	1110				DR		PCOffset9									
NOT <sup>+</sup>	1001				DR			SR		111111						
STR	0111				SR			BaseR		offset6						
HALT	1111				000000000000											

Figure 2: Instruction Format

For any instruction, the 4 most significant bits are the opcode which indicates which instruction is being performed.

For example, if an instruction said ADD R3, R3, R4, which means take the contents of register 3 and register 4, add them and store the value back into register 3, the machine code would be as follows:

0001 011 011 000 100

Some hints: Values indicated as offset or imm simply need to sign extended to the correct number of bits (such as 3 for imm5 would be written as 00011). Values called PCOffset are determined based on the value of the PC at the time. The excel sheet you're given indicates the PC at each instruction to assist you with these calculations.

### 2.2.3 Loading the instructions to the LC-M

1. Convert the instructions to machine code and place in the binary column of the excel document `simple_code` (A few of these instructions are already done for you).
2. Once complete, the excel document will have converted them into hexadecimal. Copy this hexadecimal into a blank text file called `machinecode.txt`.
3. In the main circuit, click the RST button and then right click on the RAM component and choose "Edit Contents".
4. A new window should pop up and in the upper right hand corner you should see a button that says "Load from file", which you can click on and then navigate to where you saved `machinecode.txt` and load it.

The LC-M is now set up to run a program!

## 2.3 Part 3

For this part of the assignment you will trace through the execution of the machine code in the LC-M and complete a Canvas quiz about the execution process. Feel free to add Probes to the circuit so you can keep track of the values in wires at each time. Notice that you can check the the number of clock cycles since the beginning of the execution using the counter circuit next to the main LC-M circuit.

**Note: When the question asks to view the internal state of a circuit do not look at that subcircuit directly. Instead, right-click the circuit in the LC-M and press "View Internal State"**

### 2.3.1 Information about State Machines

An important part of running a program is the state machine. The state machine keeps track of what state the computer is in (such as FETCH), and outputs the proper control signals for each state (such as the value of a selector for a multiplexer). The state machine has already been built for you so that you will be able to easily run a program, but try and pay attention to see what control signals (indicated by LEDs in the main circuit) light up during different times of a program. An important note to this is that while many instructions (such as ADD and NOT) can be implemented in one clock cycle, some (such as LDR) can take even more! For this part of the assignment you will trace through the execution of the machine code in the LC-M, and will answer the a quiz about the execution process on Canvas.

Hint: The state machine always follows this format:

FETCH → DECODE → EXECUTE

Fetch is 3 cycles that involves incrementing the PC and loading the instruction from memory into the instruction register. The decode state is 1 cycle and uses the opcode from the instruction to determine which instruction is being executed. The instruction state takes a variable number of cycles dependent on the instruction and once completed, the computer returns to the fetch state.

### 3 Testing your work

You can run the autograder using Gradescope which will test individual circuits as well as your machine code. You can also run your machine code on your completed computer! You can run the program manually by clicking the CLK on and off, but you can also use the simulator tool in Circuit Sim to run it faster. To run the simulator select the "Simulator" pulldown menu on the top of the page and make sure that "Simulation Enabled" is selected; then click on "Frequency" under the same menu to choose the clock frequency of the simulator. Finally, select "Clock Enabled" to run the simulation.

When the program ends, the isHalt LED in the main circuit will be red. When running it on the highest frequency, the program should not take more than a couple seconds to run. Notice that once you reset the simulation you need to re-load the machine code into memory to run the program again. One way to indicate that the program ran correctly is to look at the memory contents and check if the value at memory address 0xB is 0x0009. You can also check the value of each register on the Register File (DPRF) after the program ran. A correctly run program will have these values in the registers:

```
R0: 0b000000
R1: 0b01011
R2: 0b00100
R3: 0b000000
R4: 0b11101
R5: 0b000000
R6: 0b000000
R7: 0b000000
```

This isn't a guarantee that the circuit is fully correct, but a good way to measure a correctly built circuit.

### 4 Deliverables

Please upload the following files to Gradescope:

1. LCM.sim
2. machinecode.txt

Also submit the "HW5 Code Tracing" quiz on Canvas.

### 5 Demos

**This homework will be demoed.**

The demos will be ten minutes long and will occur in the CS2110 TA lab - COC 104b during the week of September 30th - October 4th.

- Demo signups will go up by September 27th and you can sign up for a demo via the calendar on Canvas
- You must sign up for your demo at least 24 hours in advance
- If you cannot make any of the available demo signups, email the Head TA at [viszlai@gatech.edu](mailto:viszlai@gatech.edu) by **Monday September 30th**. This is the only way you can ensure you will have a demo.



- Your demo is 30 points of your Homework 05 grade. If you miss your demo you will not receive these points and the maximum you can receive on the homework is a 70%.
- You cannot cancel your demo within 24 hours or else you will lose all 30 of your demo points.
- You will be able to makeup one of your demos at the end of the semester for 50% of the demo grade.

## 6 Rules and Regulations

### 6.1 General Rules

1. Starting with the assembly homeworks, any code you write must be meaningfully commented. You should comment your code in terms of the algorithm you are implementing; we all know what each line of code does.
2. Although you may ask TAs for clarification, you are ultimately responsible for what you submit. This means that (in the case of demos) you should come prepared to explain to the TA how any piece of code you submitted works, even if you copied it from the book or read about it on the internet.
3. Please read the assignment in its entirety before asking questions.
4. Please start assignments early, and ask for help early. Do not email us the night the assignment is due with questions.
5. If you find any problems with the assignment it would be greatly appreciated if you reported them to the author (which can be found at the top of the assignment). Announcements will be posted if the assignment changes.

### 6.2 Submission Conventions

1. All files you submit for assignments in this course should have your name at the top of the file as a comment for any source code file, and somewhere in the file, near the top, for other files unless otherwise noted.
2. When preparing your submission you may either submit the files individually to Canvas/Gradescope or you may submit an archive (zip or tar.gz only please) of the files. You can create an archive by right clicking on files and selecting the appropriate compress option on your system. Both ways (uploading raw files or an archive) are exactly equivalent, so choose whichever is most convenient for you.
3. Do not submit compiled files, that is .class files for Java code and .o files for C code. Only submit the files we ask for in the assignment.
4. Do not submit links to files. The autograder does not understand it, and we will not manually grade assignments submitted this way as it is easy to change the files after the submission period ends.

### 6.3 Submission Guidelines

1. You are responsible for turning in assignments on time. This includes allowing for unforeseen circumstances. If you have an emergency let us know **IN ADVANCE** of the due time supplying documentation (i.e. note from the dean, doctor's note, etc). Extensions will only be granted to those who contact us in advance of the deadline and no extensions will be made after the due date.

2. You are also responsible for ensuring that what you turned in is what you meant to turn in. After submitting you should be sure to download your submission into a brand new folder and test if it works. No excuses if you submit the wrong files, what you turn in is what we grade. In addition, your assignment must be turned in via Canvas/Gradescope. Under no circumstances whatsoever we will accept any email submission of an assignment. Note: if you were granted an extension you will still turn in the assignment over Canvas/Gradescope.
3. There is a 6-hour grace period added to all assignments. You may submit your assignment without penalty up until 11:55PM, or with 25% penalty up until 5:55AM. So what you should take from this is not to start assignments on the last day and plan to submit right at 11:54AM. You alone are responsible for submitting your homework before the grace period begins or ends; neither Canvas/Gradescope, nor your flaky internet are to blame if you are unable to submit because you banked on your computer working up until 11:54PM. The penalty for submitting during the grace period (25%) or after (no credit) is non-negotiable.

## 6.4 Syllabus Excerpt on Academic Misconduct

Academic misconduct is taken very seriously in this class. Quizzes, timed labs and the final examination are individual work.

Homework assignments are collaborative, In addition many if not all homework assignments will be evaluated via demo or code review. During this evaluation, you will be expected to be able to explain every aspect of your submission. Homework assignments will also be examined using computer programs to find evidence of unauthorized collaboration.

What is unauthorized collaboration? Each individual programming assignment should be coded by you. You may work with others, but each student should be turning in their own version of the assignment. Submissions that are essentially identical will receive a zero and will be sent to the Dean of Students' Office of Academic Integrity. Submissions that are copies that have been superficially modified to conceal that they are copies are also considered unauthorized collaboration.

**You are expressly forbidden to supply a copy of your homework to another student via electronic means. This includes simply e-mailing it to them so they can look at it. If you supply an electronic copy of your homework to another student and they are charged with copying, you will also be charged. This includes storing your code on any site which would allow other parties to obtain your code such as but not limited to public repositories (Github), pastebin, etc. If you would like to use version control, use [github.gatech.edu](http://github.gatech.edu)**

## 6.5 Is collaboration allowed?

Collaboration is allowed on a high level, meaning that you may discuss design points and concepts relevant to the homework with your peers, share algorithms and pseudo-code, as well as help each other debug code. What you shouldn't be doing, however, is pair programming where you collaborate with each other on a single instance of the code. Furthermore, sending an electronic copy of your homework to another student for them to look at and figure out what is wrong with their code is not an acceptable way to help them, because it is frequently the case that the recipient will simply modify the code and submit it as their own. Consider instead using a screen-sharing collaboration app, such as <http://webex.gatech.edu/>, to help someone with debugging if you're not in the same room.




Figure 3: Collaboration rules, explained colorfully